

Machine Learning for Credit Score Classification:

A Comparative Analysis of Ensemble Methods

Davide De Blasio
Student ID: 2082600

November 2025

1 Introduction

Credit scoring is fundamental to modern finance, enabling institutions to assess risk, determine loan eligibility, and set interest rates. Traditional methods such as FICO scores rely on limited variables and rigid rules that may not capture the full complexity of financial behavior. Machine learning offers a data-driven alternative that identifies subtle patterns and non-linear relationships across multiple behavioral and financial features.

1.1 Motivation and Objectives

This project develops and evaluates a comprehensive machine learning pipeline to predict customer credit scores (**Poor**, **Standard**, **Good**) using a Kaggle Credit Score Classification Dataset containing 100,000 records with 28 features. We compare six algorithms: Gaussian Naive Bayes, Multinomial Naive Bayes, Softmax Regression, Decision Tree, and two Random Forest variants (20 and 15 features). All models are optimized using 5-fold cross-validation with weighted F1-score as the primary metric.

The analysis focuses on four key dimensions: (1) predictive accuracy across multiple metrics, (2) model interpretability through feature importance analysis, (3) computational efficiency for practical deployment, and (4) generalization capability to unseen data. This comprehensive evaluation provides actionable insights for financial institutions seeking to implement machine learning-based credit scoring systems.

2 Dataset Description

The dataset used in this project consists of 100,000 customer records and originally contained 28 features including the target variable. The data represents diverse credit profiles from a publicly available Kaggle competition focused on credit score classification.

2.1 Features and Target Variable

The target variable, `Credit_Score`, is categorical with three balanced classes: **Poor**, **Standard**, and **Good**, representing the customer's creditworthiness level. For modeling purposes, these classes were numerically encoded as 0, 1, and 2, respectively.

The features can be categorized as follows:

- **Demographic:** Age, Occupation.
- **Financial:** Annual_Income, Monthly_Inhand_Salary, Monthly_Balance, Amount_invested_monthly, Total_EMI_per_month, Outstanding_Debt.

- **Credit Behavior:** Num_Bank_Accounts, Num_Credit_Card, Interest_Rate, Credit_Utilization_Ratio, Credit_History_Age, Credit_Mix, Changed_Credit_Limit.
- **Loan Information:** Num_of_Loan, Type_of_Loan, Delay_from_due_date, Num_of_Delayed_Payment.
- **Payment Patterns:** Payment_of_Min_Amount, Payment_Behaviour, Num_Credit_Inquiries.




Figure 1: Distribution of Target Variable (Credit Score). The dataset shows a balanced distribution across three classes: Poor, Standard, and Good, facilitating unbiased model training.

2.2 Preprocessing Pipeline

A rigorous seven-stage preprocessing pipeline was applied:

1. **Data Cleaning:** Dropped irrelevant columns (ID, Name, SSN), converted string-stored numerical features to float64, replaced negative values and placeholders with NaN.
2. **Missing Value Imputation:** Customer-specific mode/median imputation with global fallback; rows with missing Type_of_Loan deleted.
3. **Feature Engineering:** Expanded Type_of_Loan into 8 binary columns, converted Credit_History_Age to total months, re-encoded Customer_ID to integers.
4. **Statistical Feature Selection:** Chi-Square tests removed Month and Occupation. ANOVA F-tests eliminated three low-variance features (Monthly_Balance, Age, Total_EMI_per_month).
5. **Outlier Treatment:** IQR-based capping (from $Q_1 - 1.5 \times IQR$ to $Q_3 + 1.5 \times IQR$) preserved sample size while mitigating extreme values.
6. **Feature Encoding:** Target encoded as integers (Poor=0, Standard=1, Good=2), Credit_Mix ordinally encoded, Payment_Behaviour one-hot encoded with `drop_first=True`.
7. **Final Selection and Scaling:** SelectKBest with mutual information selected top 20 features from approximately 35 features. We applied an 80-20 stratified split, using RobustScaler for most models and MinMaxScaler for Multinomial NB.

2.3 Data Quality Impact

The preprocessing pipeline significantly improved data quality and feature distributions. Figure 2 shows the distribution comparison before and after preprocessing for key features.




Figure 2: Distribution Comparison: After Preprocessing. The preprocessing pipeline effectively normalized distributions, removed outliers, and improved feature quality.

3 Methodology & Models

We implemented and tuned six distinct models to solve the classification task. A 5-fold Cross-Validation strategy with `GridSearchCV` was employed for hyperparameter tuning, optimizing for the weighted F1-score.

Figure 3 shows the correlation structure among the final 20 selected features, revealing key relationships that inform model performance.

3.1 Models Implemented

1. **Gaussian Naive Bayes:** A probabilistic classifier assuming continuous features follow a Gaussian distribution.
 - *Hyperparameters:* `var_smoothing` (10^{-9} to 10^{-3}).
 - *Scaling:* `RobustScaler`.
2. **Multinomial Naive Bayes:** Suited for discrete counts.
 - *Hyperparameters:* `alpha` (0.1 to 10.0).
 - *Scaling:* `MinMaxScaler` (required for non-negative input).
3. **Softmax Regression (Logistic Regression):** A linear model generalized for multi-class problems.




Figure 3: Correlation Matrix of Top 20 Selected Features. Darker colors indicate stronger correlations, revealing feature relationships that influence model behavior.

- *Hyperparameters:* Regularization strength C (0.001 to 1.0).
- *Settings:* L2 penalty, `class_weight='balanced'` to handle class imbalance.
- *Scaling:* RobustScaler.

4. Decision Tree: A non-linear model capable of capturing complex decision boundaries.

- *Hyperparameters:* `max_depth` (5, 10, 15, None), `min_samples_split`, `min_samples_leaf`, `criterion` (Gini, Entropy).
- *Settings:* `class_weight='balanced'`.

5. Random Forest: An ensemble of decision trees to reduce variance and overfitting.

- *Hyperparameters:* `n_estimators` (150, 200), `max_depth` (20, 25), `min_samples_split`, `min_samples_leaf`.
- *Settings:* `class_weight='balanced_subsample'`.

6. Random Forest (Reduced Features): A variant of the above model using only the top 15 features (explaining 95% of cumulative importance) to test efficiency vs. performance trade-offs.

4 Performance Metrics

To evaluate the models comprehensively, we utilized the following metrics:

- **Accuracy:** The ratio of correctly predicted observations to total observations.
- **Precision, Recall, F1-Score:** Weighted averages were calculated to account for class imbalance.
- **Confusion Matrix:** To visualize misclassifications between classes.
- **ROC-AUC:** One-vs-Rest approach to measure the ability of the classifier to distinguish between classes.
- **Computational Cost:** Training time measured in seconds.

4.1 Summary of Results

Table 1 summarizes the performance of all models on the test set.

Table 1: Model Performance Comparison (Test Set)

Model	Accuracy	Precision	Recall	F1 Score
Gaussian NB	0.6174	0.6826	0.6174	0.6197
Multinomial NB	0.5780	0.5714	0.5780	0.5598
Softmax Regression	0.6535	0.6917	0.6535	0.6570
Decision Tree	0.7359	0.7560	0.7359	0.7368
Random Forest (20 Feat.)	0.7920	0.8051	0.7920	0.7928
Random Forest (15 Feat.)	0.7886	0.8026	0.7886	0.7897

Table 2: Training Performance and Computational Cost

Model	Train F1	Test F1	Gap	Time (s)
Gaussian NB	0.6174	0.6197	-0.0024	0.09
Multinomial NB	0.5611	0.5598	0.0013	0.01
Softmax Regression	0.6580	0.6570	0.0010	0.43
Decision Tree	0.8187	0.7368	0.0819	1.51
Random Forest (20 Feat.)	0.8583	0.7928	0.0655	28.56
Random Forest (15 Feat.)	0.8553	0.7897	0.0657	28.13

Table 3: Per-Class Performance for Random Forest (20 Features) - Best Model

Class	Precision	Recall	F1 Score	Support
Poor	0.77	0.87	0.82	5,533
Standard	0.87	0.73	0.80	9,309
Good	0.66	0.84	0.74	2,877
Weighted Avg.	0.81	0.79	0.79	17,719

Table 4: ROC-AUC Scores (One-vs-Rest) for All Models

Model	Poor vs Rest	Standard vs Rest	Good vs Rest
Gaussian NB	0.762	0.694	0.855
Multinomial NB	0.752	0.676	0.845
Softmax Regression	0.794	0.738	0.872
Decision Tree	0.896	0.819	0.910
Random Forest (20 Feat.)	0.933	0.869	0.956
Random Forest (15 Feat.)	0.930	0.867	0.954




Figure 4: Comparison of Accuracy & F1 Scores and Training Time across all models.



Figure 5: Comprehensive comparison of Random Forest variants (20 vs 15 features) across all performance metrics.



Figure 6: Precision-Recall histograms for all models showing trade-offs between precision and recall.




Figure 7: ROC curves (One-vs-Rest) for Random Forest showing discrimination ability for each class.

5 Results & Analysis

5.1 Model Performance and Comparative Analysis

5.1.1 Performance Hierarchy

Random Forest (20 features) achieved the best performance with $F1=0.7928$ and accuracy=0.7920, substantially outperforming all alternatives:

1. Random Forest (20 feat.): $F1 = 0.7928$
2. Random Forest (15 feat.): $F1 = 0.7897$ (0.31% lower)
3. Decision Tree: $F1 = 0.7368$ (5.60% lower)
4. Softmax Regression: $F1 = 0.6570$ (13.58% lower)
5. Gaussian NB: $F1 = 0.6197$ (17.31% lower)
6. Multinomial NB: $F1 = 0.5598$ (23.30% lower)

5.1.2 Why Random Forest Excels

Random Forest’s superiority stems from four key advantages:

1. **Ensemble learning:** 200 diverse trees reduce variance through democratic voting, improving single Decision Tree performance by 5.6%.
2. **Non-linear pattern recognition:** Recursive partitioning captures complex relationships (such as U-shaped credit utilization curves) that linear models miss.
3. **Feature interaction modeling:** Tree paths implicitly model interactions like “high debt AND low income implies poor credit”.
4. **Noise robustness:** Bagging averages out measurement errors across bootstrap samples.

The 15-feature variant achieves F1=0.7897 (only 0.31% lower), demonstrating that 25% of features contribute negligible value and can be removed for simpler deployment.

5.1.3 Impact of Model Assumptions

Naive Bayes (Gaussian F1=0.6197, Multinomial F1=0.5598) underperformed due to violated assumptions. First, feature independence is violated because credit features exhibit strong correlations, such as Annual Income with Monthly Salary, and Outstanding Debt with EMI payments. This violates the independence assumption $P(X_1, ..., X_n|Y) = \prod P(X_i|Y)$. Second, distributional assumptions fail because financial data shows heavy tails and multimodality rather than Gaussian distributions. Multinomial NB assumes discrete counts, which is inappropriate for continuous variables.

Softmax Regression (F1=0.6570) suffers from linear boundary limitations. Credit relationships are non-linear. For example, credit utilization shows U-shaped risk where both 0% and 100% are problematic, and feature interactions such as high debt combined with low income cannot be captured without explicit engineering. However, 65.7% F1 indicates some linear structure exists, justifying its use as a fast baseline.

Tree-Based Models excel due to assumption-free learning. They have no distributional requirements, automatically capture non-linearity through recursive partitioning, implicitly model interactions via tree paths, and seamlessly handle mixed data types.

5.1.4 Overfitting Analysis

Table 5 shows train-test F1 gaps quantifying overfitting. We classify models with gaps exceeding 0.10 as overfitting.


Table 5: Overfitting Analysis (Overfitting threshold: gap $\geq$ 0.10)

Model	Train F1	Test F1	Gap	Status
Gaussian NB	0.6174	0.6197	-0.0024	Underfitting
Multinomial NB	0.5611	0.5598	0.0013	Underfitting
Softmax Regression	0.6580	0.6570	0.0010	Excellent
Decision Tree	0.8187	0.7368	0.0819	Good
Random Forest (20)	0.8583	0.7928	0.0655	Good
Random Forest (15)	0.8553	0.7897	0.0657	Good


Key findings from overfitting analysis:

- **Decision Tree** shows a gap of 0.0819 (below the 0.10 threshold) by growing deep trees that memorize some training noise, but remains within acceptable bounds.
- **Random Forest** reduces the gap to 0.0655 through three mechanisms:
 1. Bagging uses 200 bootstrap samples that expose trees to different data, averaging out noise
 2. Feature randomness considers only approximately 4.5 features (square root of 20) per split, decorrelating trees
 3. Balanced class weights adjust for imbalance
- **Naive Bayes** models underfit (negative gaps) because they are too rigid to learn training data effectively.
- **Softmax Regression** shows a near-zero gap, but this reflects underfitting rather than excellent generalization.

Learning curves (Figure 8) show that Random Forest’s train-validation gap narrows from 0.10 (with 10% of data) to 0.07 (with 100% of data), indicating that more data reduces overfitting. The validation score plateaus at 0.79, suggesting capacity limits. GridSearchCV-optimized hyperparameters (`max_depth=25`, `min_samples_split=12`, `min_samples_leaf=4`, `n_estimators=200`) provide regularization.



(a) Decision Tree Learning Curve



(b) Random Forest Learning Curve

Figure 8: Learning curves showing Training vs. Validation scores.

5.1.5 Feature Importance

Figure 9 shows Random Forest feature importance (mean decrease in Gini impurity).




Figure 9: Top 20 Features by Importance.

The top 3 features account for 40.3% of total importance:

- **Credit_Mix (14.83%)**: Portfolio diversity reflects financial maturity
- **Outstanding_Debt (14.40%)**: Debt burden limits repayment capacity
- **Interest_Rate (11.07%)**: Encodes historical credit performance

Mid-tier features (ranks 4-10, accounting for 27.4%) capture payment behavior:

- Delay_from_due_date: 6.90%
- Credit_History_Age: 6.76%
- Num_Credit_Inquiries: 5.86%
- Changed_Credit_Limit: 5.54%
- Monthly_Inhand_Salary: 4.30% (surprisingly ranks only 8th, suggesting behavior matters more than income)

Key insights:

- Loan type indicators (ranks 16-20) contribute less than 0.6% each
- The top 15 features explain 95% of cumulative importance, validating the 15-feature model (F1=0.7897 versus 0.7928)
- Age and Occupation were dropped during preprocessing, suggesting that *what you do with credit* matters more than *who you are*, potentially reducing demographic bias

5.1.6 Class-Specific Performance

Table 3 and Figure 10 show per-class metrics for Random Forest.




Figure 10: Confusion Matrix for Random Forest.

Poor Class (F1=0.82):

- Highest F1 score with precision=0.77 and recall=0.87
- High recall captures 87% of risky customers
- 23% false positive rate: 1,423 Standard customers misclassified as Poor, denying credit to eligible applicants

Standard Class (F1=0.80):

- Highest precision (0.87) but lowest recall (0.73)
- Misses 27% of true Standard customers: 1,423 misclassified as Poor, 1,060 as Good
- As the majority class (52.5% of data), the model favors precision over recall

Good Class (F1=0.74):

- Lowest precision (0.66) due to minority status (16.2% of data)
- Captures 84% of good credit customers but produces 34% false positives
- High AUC (0.956) suggests good ranking ability; threshold tuning could improve precision

Key Error Patterns:

1. 1,423 Standard customers misclassified as Poor → denies credit to eligible applicants
2. 1,060 Standard customers misclassified as Good → approves moderate-risk customers at favorable terms
3. 561 Poor customers misclassified as Standard → exposes lenders to default risk

ROC-AUC Analysis (Figure 7):

- Poor vs Rest: 0.933 (excellent discrimination)
- Standard vs Rest: 0.869 (weakest due to overlap with both extremes)

- Good vs Rest: 0.956 (best discrimination)
- High AUC with low precision for Good class indicates threshold calibration issue rather than ranking problem

Figure 6 shows the precision-recall trade-offs for all models, highlighting that Random Forest achieves superior balance across all operating points.

5.1.7 Computational Cost

Training times (Table 2) vary by a factor of 3,000:

- **Fastest:** Naive Bayes (9-86 milliseconds)
- **Moderate:** Softmax Regression (0.429 seconds) - best speed-accuracy compromise
- **Slowest:** Random Forest (28.56 seconds for 20 features, 28.13 seconds for 15 features)
 - Trains 200 trees $\times$ 5 CV folds = 1,000 tree trainings
 - Acceptable for monthly retraining schedules

Deployment Recommendations:

1. Batch Processing:

- Use Random Forest (20 features) to maximize accuracy
- Training overhead negligible for monthly/quarterly updates

2. Real-Time APIs:

- Option 1: Random Forest (15 features) - 28.13s training, 1-5ms inference
- Option 2: Softmax Regression - 0.43s training, F1=0.6570
- RF inference supports 500 predictions/second on single CPU core

3. High-Frequency Retraining:

- Use Softmax Regression for sub-second retraining cycles

Performance Optimization:

- **Parallelization:** 16-core CPU reduces RF training to 3-5 seconds
- **GPU Acceleration:** Achieves 10-100 $\times$ speedup
- **Inference Time:** RF (0.5-2ms/sample) vs Softmax (0.005ms/sample)

Figure 5 provides a comprehensive comparison between the 20-feature and 15-feature Random Forest variants, showing minimal performance degradation with 25% feature reduction.

6 Conclusion

This study successfully developed and evaluated a comprehensive machine learning pipeline for multi-class credit score classification. Through rigorous preprocessing and systematic comparison of six algorithms, we transformed a raw dataset of 100,000 customer records into actionable predictive models suitable for deployment in financial institutions.

6.1 Key Findings

Best Model Performance:

- Random Forest (20 features) achieved F1=0.7928, Accuracy=0.7920
- Train-test gap=0.0655 (below 0.10 threshold, indicating good generalization)
- ROC-AUC scores: Poor (0.933), Standard (0.869), Good (0.956)

Feature Importance Insights:

- Top 3 features account for 40.3% of importance:
 1. Credit_Mix: 14.83%
 2. Outstanding_Debt: 14.40%
 3. Interest_Rate: 11.07%
- Top 15 features explain 95% of importance
- 15-feature Random Forest achieves F1=0.7897 (only 0.31% lower)

Model Assumptions Impact:

- Naive Bayes underperformed by 17-23% due to feature independence violations
- Softmax Regression achieved 13.58% lower F1 due to linear boundary limitations
- Tree-based models' assumption-free learning proved optimal for heterogeneous credit data

Overfitting Control:

- Decision Tree gap: 0.0819 (below 0.10 threshold)
- Random Forest reduced gap to 0.0655 through:
 1. Bagging across 200 bootstrap samples
 2. Feature randomness (approximately 4.5 features per split)
 3. Balanced class weights
- GridSearchCV hyperparameters provide effective regularization

6.2 Practical Implications

Deployment Recommendations by Use Case:

1. Batch Processing:

- Use Random Forest with 20 features to maximize accuracy
- Monthly/quarterly retraining schedules easily accommodate 28.56-second training time

2. Real-Time Applications:

- Option 1: Random Forest (15 features) - 1-5ms inference
- Option 2: Softmax Regression - 0.43s training for frequent updates

3. Explainability Requirements:

- Focus on top 5 features (account for over 50% of decisions)
- Satisfies regulatory transparency requirements

Class-Specific Improvement Strategies:

- **Poor Class (F1=0.82):**
 - Reduce 23% false positive rate
 - Avoid denying credit to 1,423 misclassified Standard customers
 - Implement secondary review process for borderline cases
- **Standard Class (F1=0.80):**
 - Improve 73% recall through enhanced features
 - Implement secondary manual reviews for edge cases
 - Target the 2,483 misclassified customers
- **Good Class (F1=0.74):**
 - Improve 66% precision using cost-sensitive learning
 - Apply threshold tuning (high AUC=0.956 indicates good ranking)
 - Consider collecting more data for minority class (16.2% of dataset)

6.3 Limitations

Data Limitations:

- Missing key contextual features: employment history, education level
- Lacks temporal trends: cross-sectional data ignores credit behavior evolution
- No macroeconomic variables: unemployment rate, inflation, interest rate environment
- Geographic information absent: models may not generalize across regions

Modeling Assumptions:

- Assumes Kaggle labels accurately reflect true credit risk
- IID assumption violated: correlated economic shocks affect multiple customers simultaneously
- Static model: does not account for concept drift or changing credit standards

6.4 Future Work

Advanced Algorithms:

- Evaluate gradient boosting methods (XGBoost, LightGBM, CatBoost)
- Expected 2-5% F1 improvement through sequential error correction
- Temporal models (LSTMs) to capture credit trajectory trends over time

Enhanced Feature Engineering:

- Polynomial interactions: $Outstanding_Debt \times Interest_Rate$

- Derived ratios: debt-to-income, credit utilization trends
- Temporal features: payment consistency, credit mix evolution

Class Imbalance Techniques:

- Note: SMOTE and SMOTE+Tomek were tested but yielded worse performance than balanced class weights
- Explore cost-sensitive learning with asymmetric misclassification penalties
- Threshold optimization per class using precision-recall curves

Model Validation:

- Fairness auditing to ensure equitable lending practices
- Adversarial testing for robustness
- External validation on different credit score datasets

6.5 Concluding Remarks

This study demonstrates that machine learning can significantly enhance credit scoring accuracy while maintaining interpretability:

- **Production-Ready Performance:** Random Forest achieved 79.2% accuracy (F1=0.7928) with good generalization (train-test gap=0.0655, below the 0.10 overfitting threshold).
- **Behavior Over Demographics:** Behavioral features (Credit_Mix, Outstanding_Debt, Payment Delays) vastly outweigh demographics (Age and Occupation were dropped), suggesting a paradigm shift toward behavior-centric scoring that potentially reduces historical biases.
- **Interpretable ML:** Feature importance analysis and ROC-AUC curves provide the transparency required in regulated financial contexts, enabling explainable credit decisions.
- **Practical Deployment:** The model supports multiple deployment scenarios from batch processing to real-time APIs, with inference times of 1-5 milliseconds per prediction.
- **Future Potential:** Integration of temporal modeling, advanced ensemble methods, and fairness-aware learning could push performance toward the 85-90% F1 frontier.

The convergence of rigorous preprocessing, assumption-free learning, and comprehensive evaluation positions ensemble methods as the foundation for next-generation credit risk assessment systems.